ACM Projects Team [TBD]

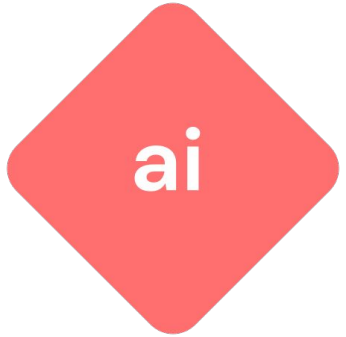

Colors!

here are the main colors we use at acm!

- binary blue
- big O(range)
- ctf cyan
- prototyping pink
- innovation indigo
- sentient scarlet

Meet the Team!

Vincent Tu
Mentor

Sia Patodia

Aryaman Dayal

Hargen Zheng

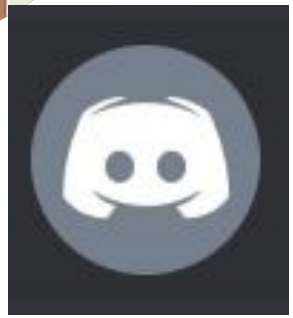

Catherine Zhang

Ryan Wong

Phillip Wu

Inspiration & Background

- AI Chatbot
  - Responds based on user input
- MBTI Personality Test
  - Can influence texting style

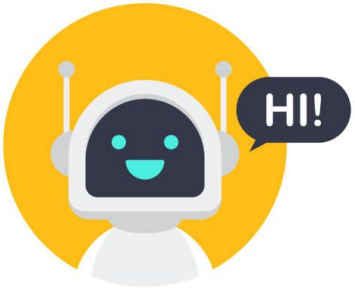

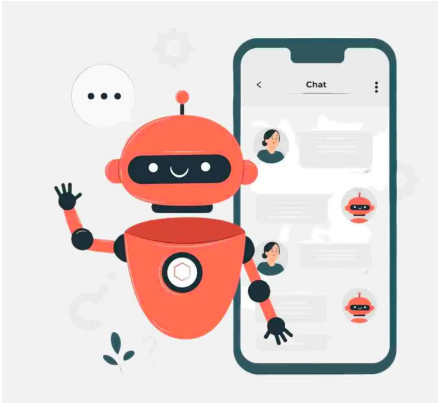

Features

- ✓ Predict MBTI personality with a decent accuracy
- ○ Have a chatbot talk back to user with that MBTI personality

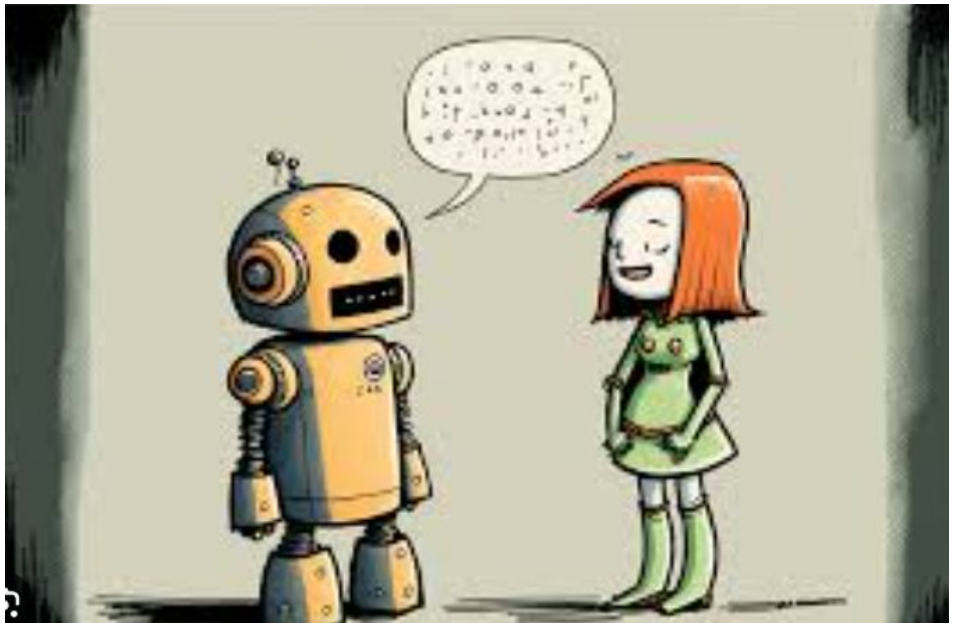

Look at our tech stack !!

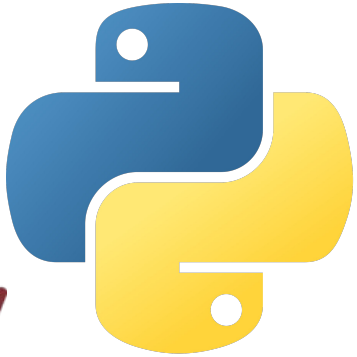

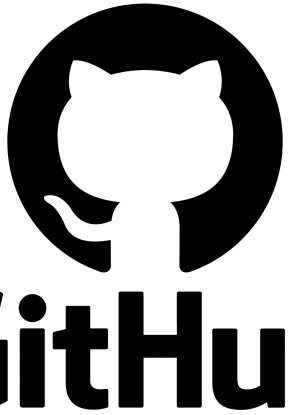

Technical Process

Dataset link: [(MBTI) Myers-Briggs Personality Type Dataset]

- EDA results:
  - Dataset has more introverts (super unbalanced)
  - Each entry has user's MBTI type and words from user's last 50 posted tweets
- Preprocessing:
  - Tokenization using BERT pre-trained tokenizer
  - Encode the MBTI type into integers 0-15
  - Only keep words between 3 and 30 characters

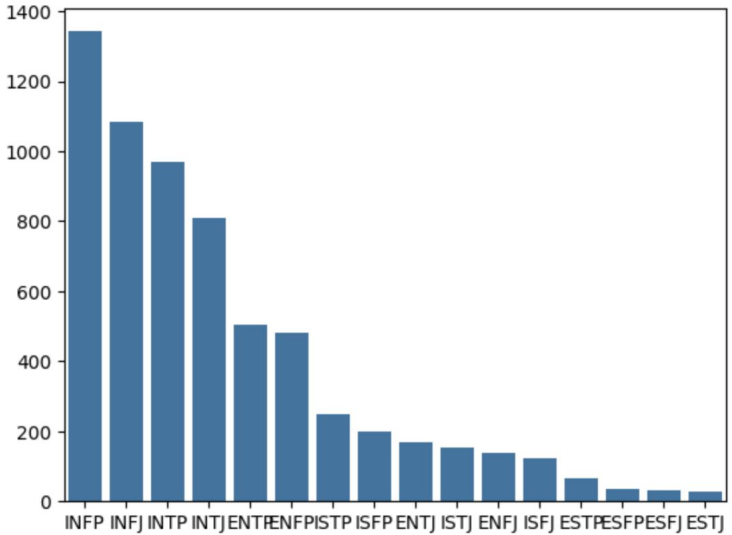

INFP	21%
INFJ	17%
Other (5373)	62%

8675
unique values

Here are the original notes vincent wrote down

- Talk about dataset (e.g. dataset size, modality, feature engineering, cleaning, EDA, insights, conclusions, motivation behind this dataset)
- Talk about preprocessing (tokenization, stemming, lemmatizing, etc)
- Talk about the model (BERT other neural language models -> talk more about how these models are made and how they work; LSTM -> long-short-term memory loosely inspired by human brain)
- Talk about the training (and finetuning), discuss training configuration and set up; if you CV, then you can also talk about that
- Talk about your experiments (what worked and what didn't work; possibly hypotheses for why things didn't work or why things worked)
- Talk inference (how do we run the model)?

Technical Process (cont.)

- Model: **BERT (Bidirectional Encoder Representations from Transformers)**
  - Unique since it is considered a bidirectional model (captures context well)
  - Uses a transformer architecture
  - Combined predicting words along with predicting if sentences belonged to each other to create the BERT model

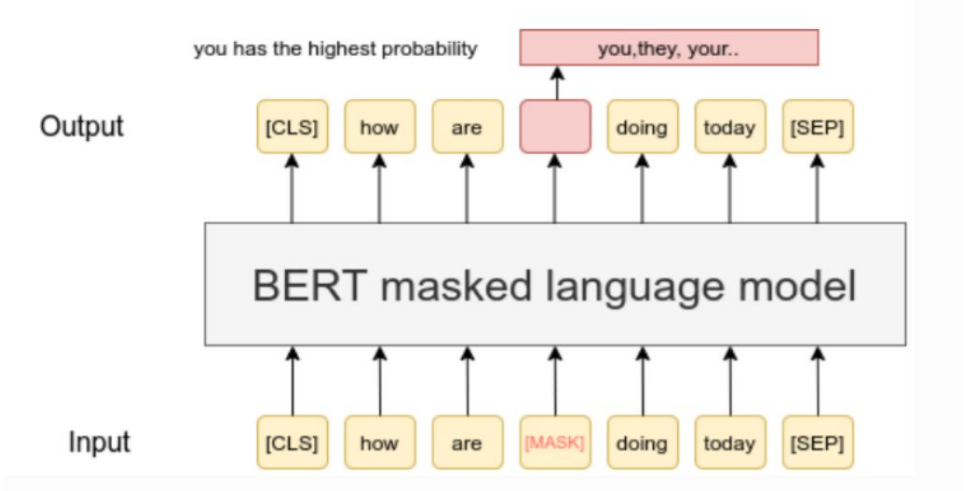

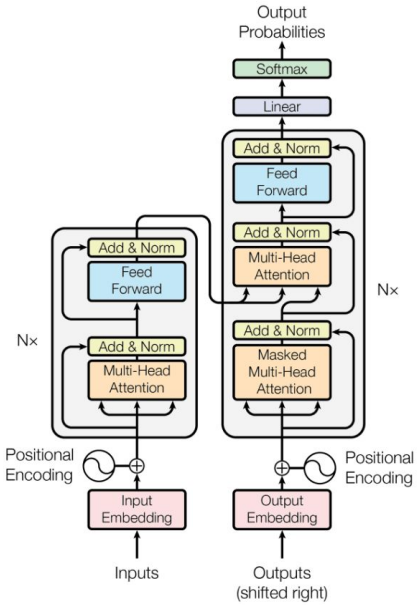

Technical Process (cont.)

Training the BERT Model

- Transfer learning (cold start issue due to size of parameters, limited computation power). We used Colab T4 GPU to train the model.
- Embed each text input into a vector with maximum length of 256 (padding).
- Shuffle the data and use 80/20 train-validation split.
- Batch size of 16 (drop the remainder). This results in 433 batches in each epoch.
- Add a fully connected layer with 512 hidden units, with ReLU activation function.
- Use a softmax layer with 16 output units to represent the probabilities of each MBTI type, given the input text corpus.
- Used the built-in Adam Optimizer with decay and Cross Entropy Loss function.

Technical Process (cont.)

BERT Model Summary

```
 Layer (type)                Output Shape              Param #   Connected to
==================================================================================================
 input_ids (InputLayer)      [(None, 256)]             0         []

 attention_mask (InputLayer  [(None, 256)]             0         []
 )

 bert (TFBertMainLayer)      TFBaseModelOutputWithPooli 108310272 ['input_ids[0][0]',
                             ngAndCrossAttentions(last_           'attention_mask[0][0]']
                             hidden_state=(None, 256, 7
                             68),
                              pooler_output=(None, 768)
                             , past_key_values=None, hi
                             dden_states=None, attentio
                             ns=None, cross_attentions=
                             None)

 intermediate_layer (Dense)  (None, 512)               393728    ['bert[0][1]']

 output_layer (Dense)        (None, 16)                8208      ['intermediate_layer[0][0]']

==================================================================================================
Total params: 108712208 (414.70 MB)
Trainable params: 108712208 (414.70 MB)
Non-trainable params: 0 (0.00 Byte)
```

Technical Process (cont.)

Experiments

- Though 1e-5 is the recommended learning rate for most of the transformer models, it was super slow to train for us. As a rookie I changed to 1e-4 and we were overshooting – after 1 epoch, the accuracy went from 35% down to 18% and it kept decreasing.
- Used 5e-5 as learning rate and decay of 1e-6 – accuracy keeps increasing all the way to ~95% on the training set. This is way better than the original Softmax regression model, which obtained a 84% on the training set.
- Could've tune the maximum size of word embeddings and the size of hidden layers

Technical Process (cont.)

Eventual Model Performance

```
Epoch 1/2
433/433 [==============================] - 423s 976ms/step - loss: 0.1813 - accuracy: 0.9433 - val_loss: 0.0714 - val_accuracy: 0.9776
Epoch 2/2
433/433 [==============================] - 414s 956ms/step - loss: 0.1713 - accuracy: 0.9479 - val_loss: 0.0697 - val_accuracy: 0.9794
```

Technical Process (cont.)

Inference

- We clean the text corpus in the same way as we did for the training and validation examples.
- Then, we used the BERT tokenizer to embed the input text corpus.
- The feature vector of the input text is then fed into our transformer model, which gives us 16 probabilities corresponding to the likelihood of each personality type.
- We then extract the top three most likely personalities and display to the user.

Technical Process (cont.)

Streamlit App

- We created a Streamlit application that shows the distribution of MBTI personality types in our dataset.
- User could input paragraphs of text and the app will predict their top three likely MBTI personalities using our model.
- WHY TOP 3, not just 1???
- After the prediction result pops up, there will be additional MBTI information of the three likely personalities, so the user can learn more about their potential MBTI type.

Challenges

- Meeting scheduling conflicts
- Balancing classwork and project
- Lack of experience (PyTorch, Model Choice, Training)
- Model overfitting

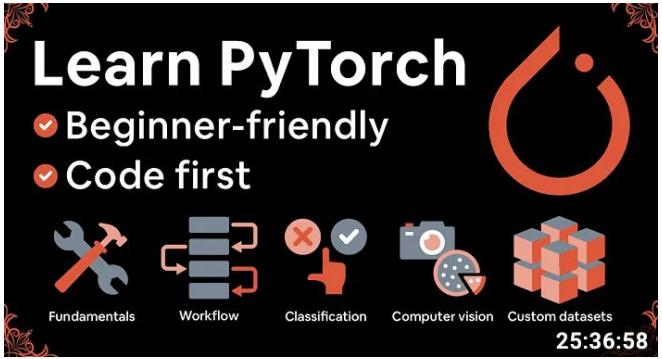

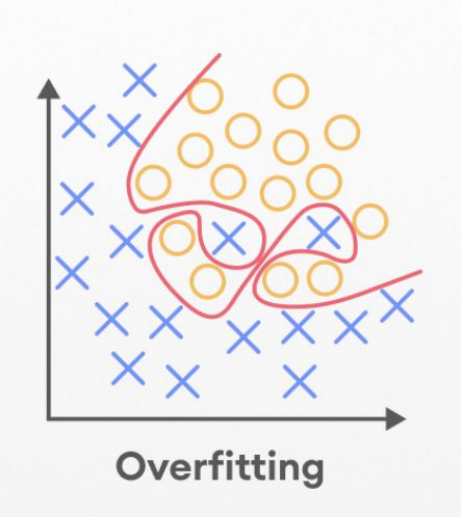

Reflection

- Smaller group sizes may make it easier to meet
- Self-guided learning is difficult
- Talk about how RoBERTa could potentially work better (as it was trained on a larger text corpus)

Where do we go from here?

- Create mobile application for interacting with chatbot
- Learn more about NLP and deep learning

Demo

Questions?

Thank you!

- Mentor: Vincent Tu LinkedIn | GitHub
- Catherine Zhang LinkedIn | GitHub
- Aryaman Dayal LinkedIn | GitHub
- Hargen Zheng LinkedIn | GitHub
- Sia Patodia LinkedIn | GitHub
- Phillip Wu GitHub
- Ryan Wong LinkedIn | GitHub